

# 1: Import Necessary Libraries and Load Dataset

```
In [1]: # Data Handling
import pandas as pd
import numpy as np

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Preprocessing
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split

# Model Evaluation
from sklearn.metrics import classification_report, confusion_matrix

# Deep Learning
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout

# Load dataset
df = pd.read_csv('Telco-Customer-Churn.csv')

# Display top rows
df.head()
```

```
Out[1]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	..
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	..
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	..
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes	..
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Yes	..
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	No	..

5 rows × 21 columns

## 2: Exploratory Data Analysis (EDA)

```
In [2]: # Dataset structure
df.info()

# Check for missing values
df.isnull().sum()

# TotalCharges has 11 blank values - investigate
df[df['TotalCharges'].isnull()]

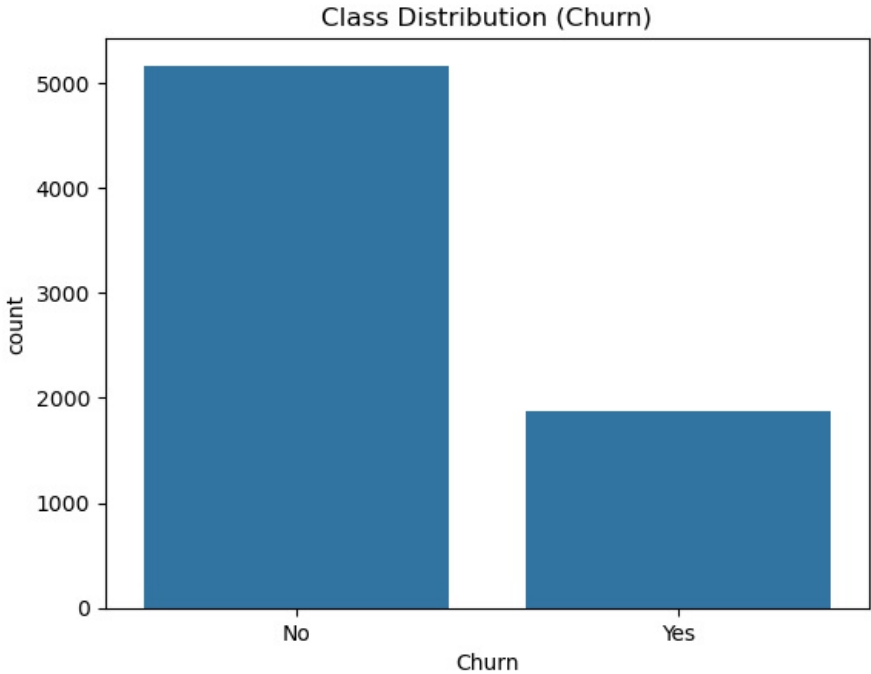
# Convert TotalCharges to numeric
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')

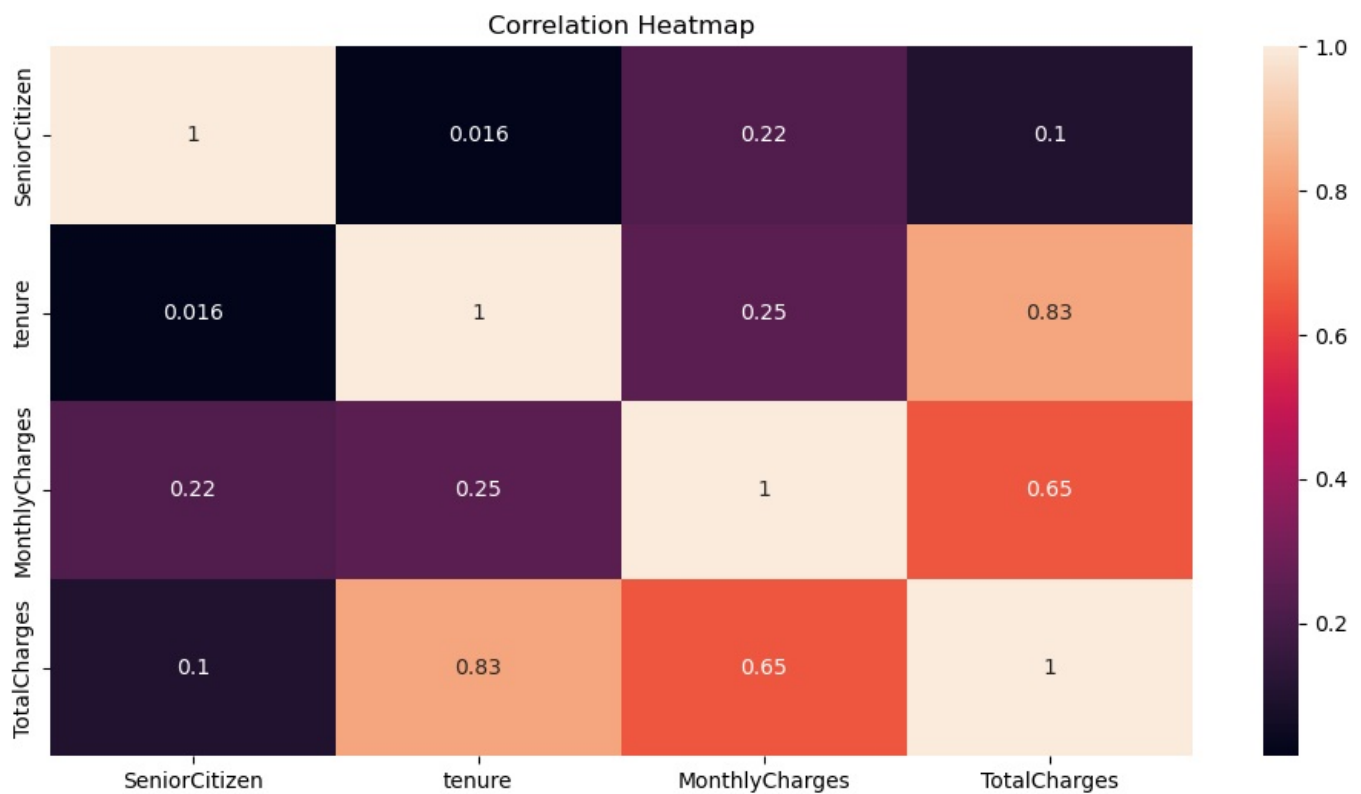
# Drop rows with missing TotalCharges
df.dropna(subset=['TotalCharges'], inplace=True)

# Check for class balance
sns.countplot(data=df, x='Churn')
plt.title("Class Distribution (Churn)")
plt.show()

# Quick correlation heatmap
plt.figure(figsize=(12, 6))
sns.heatmap(df.corr(numeric_only=True), annot=True)
plt.title("Correlation Heatmap")
plt.show()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                7043 non-null   object
2   SeniorCitizen         7043 non-null   int64
3   Partner               7043 non-null   object
4   Dependents            7043 non-null   object
5   tenure               7043 non-null   int64
6   PhoneService          7043 non-null   object
7   MultipleLines         7043 non-null   object
8   InternetService       7043 non-null   object
9   OnlineSecurity        7043 non-null   object
10  OnlineBackup          7043 non-null   object
11  DeviceProtection      7043 non-null   object
12  TechSupport           7043 non-null   object
13  StreamingTV           7043 non-null   object
14  StreamingMovies       7043 non-null   object
15  Contract              7043 non-null   object
16  PaperlessBilling      7043 non-null   object
17  PaymentMethod         7043 non-null   object
18  MonthlyCharges        7043 non-null   float64
19  TotalCharges          7043 non-null   object
20  Churn                 7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```





### 3: Feature Engineering and Encoding

```
In [6]: # Drop customerID
df.drop('customerID', axis=1, inplace=True)

# Binary features: Label Encoding
binary_cols = ['gender', 'Partner', 'Dependents', 'PhoneService',
               'PaperlessBilling', 'Churn']

le = LabelEncoder()
for col in binary_cols:
    df[col] = le.fit_transform(df[col])

# Multicategory columns: One-Hot Encoding
multi_cat_cols = ['MultipleLines', 'InternetService', 'OnlineSecurity',
                  'OnlineBackup', 'DeviceProtection', 'TechSupport',
                  'StreamingTV', 'StreamingMovies', 'Contract', 'PaymentMethod']

df = pd.get_dummies(df, columns=multi_cat_cols)

# Feature Scaling (tenure, MonthlyCharges, TotalCharges)
scaler = StandardScaler()
df[['tenure', 'MonthlyCharges', 'TotalCharges']] = scaler.fit_transform(df[['tenure', 'MonthlyCharges', 'TotalCharges']])

# Separate features and target
X = df.drop('Churn', axis=1)
y = df['Churn']
```

### 4: Train-Test Split and Validation Set

```
In [7]: # Train-test split (80/20 with validation in training set)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y, random_state=42)

# Further split training set to get validation data
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.2, stratify=y_train, random_state=42)
```

### 5: Build and Train ANN with Dropout Regularization

```
In [8]: # Define model. Here I used Dropout randomly drop 30% of Neurons from the previous layer because the model overfits
model = Sequential()
model.add(Dense(64, input_dim=X_train.shape[1], activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(32, activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(1, activation='sigmoid'))

# Compile model
```

```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

```
# Train model
history = model.fit(X_train, y_train,
                    validation_data=(X_val, y_val),
                    epochs=50, batch_size=64, verbose=1)
```

Epoch 1/50

C:\Users\jerome\anaconda3\Lib\site-packages\keras\src\layers\core\dense.py:93: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

71/71 ————— 2s 6ms/step - accuracy: 0.6479 - loss: 0.6191 - val\_accuracy: 0.7938 - val\_loss: 0.4475

Epoch 2/50

71/71 ————— 0s 3ms/step - accuracy: 0.7956 - loss: 0.4428 - val\_accuracy: 0.7893 - val\_loss: 0.4360

Epoch 3/50

71/71 ————— 0s 4ms/step - accuracy: 0.7868 - loss: 0.4554 - val\_accuracy: 0.7876 - val\_loss: 0.4361

Epoch 4/50

71/71 ————— 0s 4ms/step - accuracy: 0.7974 - loss: 0.4397 - val\_accuracy: 0.7849 - val\_loss: 0.4332

Epoch 5/50

71/71 ————— 0s 3ms/step - accuracy: 0.8075 - loss: 0.4169 - val\_accuracy: 0.7867 - val\_loss: 0.4306

Epoch 6/50

71/71 ————— 0s 3ms/step - accuracy: 0.8052 - loss: 0.4202 - val\_accuracy: 0.7840 - val\_loss: 0.4304

Epoch 7/50

71/71 ————— 0s 3ms/step - accuracy: 0.8127 - loss: 0.4122 - val\_accuracy: 0.7876 - val\_loss: 0.4303

Epoch 8/50

71/71 ————— 0s 4ms/step - accuracy: 0.8056 - loss: 0.4208 - val\_accuracy: 0.7876 - val\_loss: 0.4304

Epoch 9/50

71/71 ————— 0s 3ms/step - accuracy: 0.8055 - loss: 0.4294 - val\_accuracy: 0.7858 - val\_loss: 0.4312

Epoch 10/50

71/71 ————— 0s 3ms/step - accuracy: 0.8138 - loss: 0.4090 - val\_accuracy: 0.7893 - val\_loss: 0.4290

Epoch 11/50

71/71 ————— 0s 3ms/step - accuracy: 0.7942 - loss: 0.4202 - val\_accuracy: 0.7884 - val\_loss: 0.4298

Epoch 12/50

71/71 ————— 0s 3ms/step - accuracy: 0.8070 - loss: 0.4104 - val\_accuracy: 0.7822 - val\_loss: 0.4306

Epoch 13/50

71/71 ————— 0s 3ms/step - accuracy: 0.8091 - loss: 0.4007 - val\_accuracy: 0.7876 - val\_loss: 0.4302

Epoch 14/50

71/71 ————— 0s 3ms/step - accuracy: 0.8199 - loss: 0.4048 - val\_accuracy: 0.7884 - val\_loss: 0.4293

Epoch 15/50

71/71 ————— 0s 3ms/step - accuracy: 0.8107 - loss: 0.4063 - val\_accuracy: 0.7929 - val\_loss: 0.4310

Epoch 16/50

71/71 ————— 0s 4ms/step - accuracy: 0.8236 - loss: 0.3883 - val\_accuracy: 0.7920 - val\_loss: 0.4332

Epoch 17/50

71/71 ————— 0s 3ms/step - accuracy: 0.8141 - loss: 0.3975 - val\_accuracy: 0.7929 - val\_loss: 0.4307

Epoch 18/50

71/71 ————— 0s 3ms/step - accuracy: 0.8069 - loss: 0.4045 - val\_accuracy: 0.7938 - val\_loss: 0.4308

Epoch 19/50

71/71 ————— 0s 3ms/step - accuracy: 0.8078 - loss: 0.4035 - val\_accuracy: 0.7920 - val\_loss: 0.4328

Epoch 20/50

71/71 ————— 0s 3ms/step - accuracy: 0.8155 - loss: 0.4008 - val\_accuracy: 0.7902 - val\_loss: 0.4315

Epoch 21/50

71/71 ————— 0s 3ms/step - accuracy: 0.8109 - loss: 0.3991 - val\_accuracy: 0.7911 - val\_loss: 0.4317

Epoch 22/50

71/71 ————— 0s 3ms/step - accuracy: 0.8224 - loss: 0.3978 - val\_accuracy: 0.7929 - val\_loss: 0.4337



























Epoch 23/50

71/71 ————— 0s 3ms/step - accuracy: 0.8280 - loss: 0.3771 - val\_accuracy: 0.7929 - val\_loss: 0.4341

Epoch 24/50

71/71 ————— 0s 4ms/step - accuracy: 0.8201 - loss: 0.3927 - val\_accuracy: 0.7964 - val\_loss: 0.4349

```

Epoch 25/50
71/71  0s 4ms/step - accuracy: 0.8232 - loss: 0.3962 - val_accuracy: 0.7893 - val_loss: 0.4359
Epoch 26/50
71/71  0s 3ms/step - accuracy: 0.8161 - loss: 0.3950 - val_accuracy: 0.7938 - val_loss: 0.4353
Epoch 27/50
71/71  0s 3ms/step - accuracy: 0.8097 - loss: 0.4066 - val_accuracy: 0.7929 - val_loss: 0.4349
Epoch 28/50
71/71  0s 3ms/step - accuracy: 0.8202 - loss: 0.3944 - val_accuracy: 0.7964 - val_loss: 0.4345
Epoch 29/50
71/71  0s 3ms/step - accuracy: 0.8275 - loss: 0.3846 - val_accuracy: 0.7911 - val_loss: 0.4355
Epoch 30/50
71/71  0s 3ms/step - accuracy: 0.8219 - loss: 0.3913 - val_accuracy: 0.7956 - val_loss: 0.4353
Epoch 31/50
71/71  0s 3ms/step - accuracy: 0.8263 - loss: 0.3817 - val_accuracy: 0.7938 - val_loss: 0.4349
Epoch 32/50
71/71  0s 3ms/step - accuracy: 0.8191 - loss: 0.3937 - val_accuracy: 0.7938 - val_loss: 0.4365
Epoch 33/50
71/71  0s 3ms/step - accuracy: 0.8144 - loss: 0.3891 - val_accuracy: 0.7938 - val_loss: 0.4354
Epoch 34/50
71/71  0s 3ms/step - accuracy: 0.8224 - loss: 0.3880 - val_accuracy: 0.7884 - val_loss: 0.4372
Epoch 35/50
71/71  0s 3ms/step - accuracy: 0.8155 - loss: 0.3930 - val_accuracy: 0.7991 - val_loss: 0.4383
Epoch 36/50
71/71  0s 3ms/step - accuracy: 0.8229 - loss: 0.3900 - val_accuracy: 0.7947 - val_loss: 0.4365
Epoch 37/50
71/71  0s 3ms/step - accuracy: 0.8134 - loss: 0.3960 - val_accuracy: 0.7938 - val_loss: 0.4400
Epoch 38/50
71/71  0s 3ms/step - accuracy: 0.8165 - loss: 0.3916 - val_accuracy: 0.7920 - val_loss: 0.4383
Epoch 39/50
71/71  0s 3ms/step - accuracy: 0.8217 - loss: 0.3846 - val_accuracy: 0.7947 - val_loss: 0.4392
Epoch 40/50
71/71  0s 4ms/step - accuracy: 0.8155 - loss: 0.3893 - val_accuracy: 0.7929 - val_loss: 0.4398
Epoch 41/50
71/71  0s 3ms/step - accuracy: 0.8206 - loss: 0.3762 - val_accuracy: 0.7956 - val_loss: 0.4385
Epoch 42/50
71/71  0s 3ms/step - accuracy: 0.8190 - loss: 0.3794 - val_accuracy: 0.7920 - val_loss: 0.4417
Epoch 43/50
71/71  0s 3ms/step - accuracy: 0.8179 - loss: 0.3782 - val_accuracy: 0.7973 - val_loss: 0.4447
Epoch 44/50
71/71  0s 3ms/step - accuracy: 0.8239 - loss: 0.3774 - val_accuracy: 0.7938 - val_loss: 0.4444
Epoch 45/50
71/71  0s 3ms/step - accuracy: 0.8349 - loss: 0.3666 - val_accuracy: 0.7929 - val_loss: 0.4457
Epoch 46/50
71/71  0s 3ms/step - accuracy: 0.8232 - loss: 0.3876 - val_accuracy: 0.7938 - val_loss: 0.4447
Epoch 47/50
71/71  0s 3ms/step - accuracy: 0.8086 - loss: 0.3910 - val_accuracy: 0.7876 - val_loss: 0.4464
Epoch 48/50
71/71  0s 3ms/step - accuracy: 0.8265 - loss: 0.3805 - val_accuracy: 0.7947 - val_loss: 0.4463
Epoch 49/50
71/71  0s 3ms/step - accuracy: 0.8241 - loss: 0.3774 - val_accuracy: 0.7884 - val_loss: 0.4460
Epoch 50/50
71/71  0s 3ms/step - accuracy: 0.8222 - loss: 0.3789 - val_accuracy: 0.7956 - val_loss: 0.4466



```

## 6: Visualize Loss and Accuracy Curves

```
In [9]: # Accuracy plot
```

```
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Val Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

# Loss plot
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



## 7: Evaluate on Test Data


```
In [12]: # Predictions
y_pred_prob = model.predict(X_test)
y_pred = (y_pred_prob > 0.5).astype(int)

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
```

```
plt.ylabel("Actual")
plt.show()

# Classification Report
print(classification_report(y_test, y_pred, target_names=['No Churn', 'Churn']))
```

44/44 ————— 0s 3ms/step



	precision	recall	f1-score	support
No Churn	0.84	0.89	0.87	1033
Churn	0.64	0.53	0.58	374
accuracy	0.80			1407
macro avg	0.74	0.71	0.72	1407
weighted avg	0.79	0.80	0.79	1407

## 8: Conclusion

This project successfully developed an Artificial Neural Network (ANN) using TensorFlow and Keras to predict customer churn for a telecommunications company. The data was thoroughly preprocessed through exploratory data analysis, label and one-hot encoding, handling missing values, and scaling numerical features.

The final model included two hidden layers with dropout regularization to mitigate overfitting. Despite initial signs of overfitting during training, the model achieved a respectable 80% test accuracy. However, performance metrics showed a precision of 0.64 and recall of 0.53 for the 'Churn' class, indicating that the model struggled somewhat with identifying actual churners.

Nevertheless, the ANN effectively learned from patterns in customer attributes such as contract type, tenure, internet service, and billing method, which were key indicators of churn.

## 9: Recommendations

### 1. Improve Model Recall for Churners

Since retaining customers is usually more valuable than misclassifying non-churners, focus should be placed on increasing recall for the 'Churn' class. This can be achieved by:

- Using class weighting in the loss function to penalize false negatives.
- Exploring SMOTE or other resampling techniques to balance the training data.

### 2. Hyperparameter Tuning & Model Optimization

- Consider grid/random search or using KerasTuner to fine-tune batch size, number of neurons, and learning rate.
- Try alternate architectures (e.g., more layers, different activation functions).
- Test advanced models like XGBoost or ensemble methods to benchmark performance.

### 3. Feature Importance and Business Action

- Features like Contract type, Tech Support, Monthly Charges, and Tenure significantly influence churn.
- Prioritize retention campaigns for customers:
  - On Month-to-month contracts
  - With short tenure
  - Without technical support or online backup

### 4. Deploy and Monitor Model

- Integrate the model into the company's CRM system to flag at-risk customers.
- Retrain periodically using fresh data to maintain performance.

### 5. Customer Retention Strategy

- Introduce loyalty programs, discounts for long-term contracts, and proactive support services to reduce churn likelihood among high-risk segments.

In [ ]:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js